Agentic AI Security and AI security

Chunchi Peter Liu

AI Agent versus AI Infrastructure

AI Agent

- Definition: An application that manages **stateful information** between caller and LLM/tools, using which to complete ambiguous tasks.

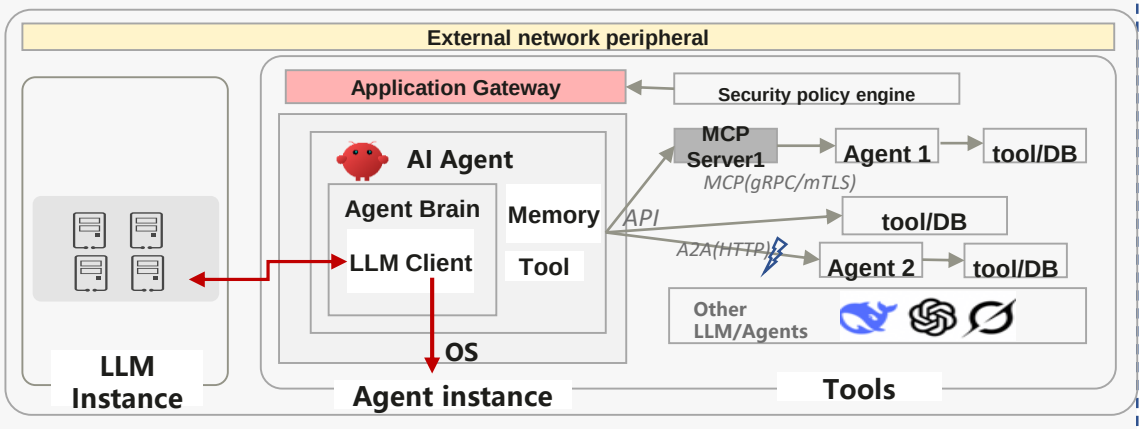

AI Agent *is* a workload, but slightly special:

1. Maintain stateful information: like credentials, identities and history contexts
2. Accept ambiguous inputs
3. More unpredictable output and resource calls

AI Infrastructure/AI-DC

- Definition: Specialized hardware, software and network stack to support the computational requirements of AI applications.

Application Layer

AI Platform Layer

Orchestration Layer

Physical Layer

Comparison	AI-DC	Regular DC
Chips	GPU/ASIC/FPGA	Mostly CPU
Apps	AI-oriented application: training and inference	Trad. enterprise IT systems, IDC, webhosting
Network	next-gen high-speed interconnect (IB/RoCE), large scale P2P arch	Standard ethernet and hierarchical architecture

Risk Analysis: OWASP published Agentic AI Application Top10, a compound result of LLM TOP10 + other GAPs

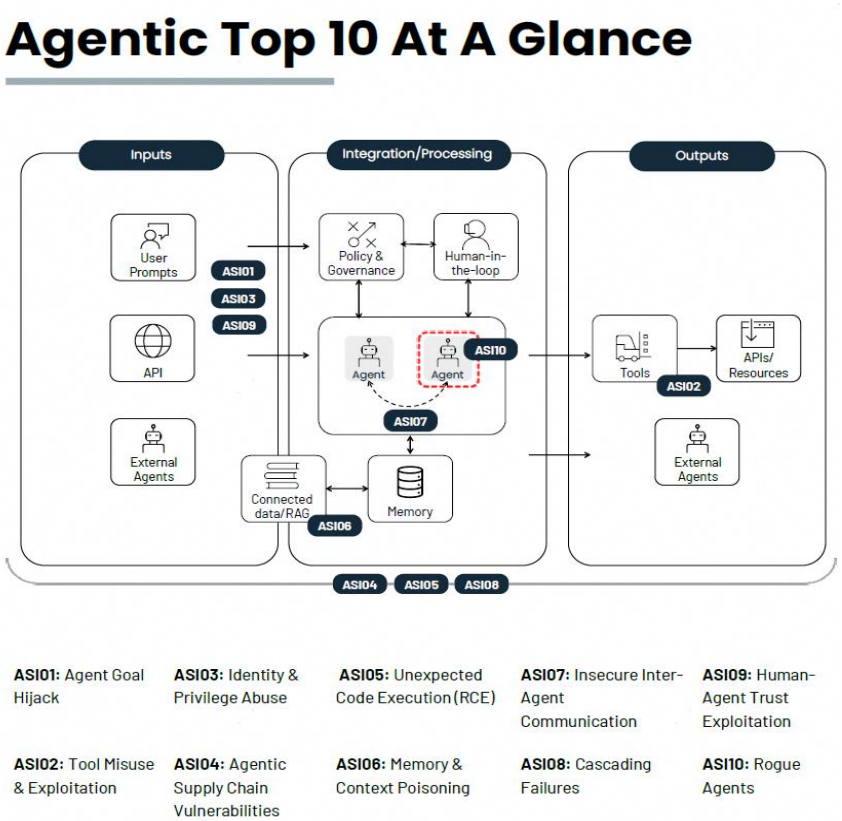

Agent Top10	LLM TOP10	Risk Analysis and other GAPs	AI Agent-specific GAPs
ASI01 Agent Goal Hijack	LLM01 prompt injection LLM06 excessive agency	Due to agents' difficulty in distinguishing legitimate instructions from attacker-controlled content,, attackers can use prompt injection (LLM01) to exploit excessive agency(LLM06), guiding agents to unintended actions.	• LLM Input/Output Guardrails • Fine-grained access control
ASI02 Tool Misuse & Exploitation	LLM06 excessive agency	Excessive agency (LLM06), combined with missing security checks when calling tools, leads to their insecure use within permissions.	• Fine-grained access control
ASI03 Identity & Privilege Abuse	LLM01 prompt injection LLM06 excessive agency LLM02 sensitive Info disclosure	Due to the lack of agent identity, arbitrary inheritance of permissions, and excessive trust among multiple agents, prompt injection and excessive agency lead to impersonation or privilege abuse， amplifying sensitive Info disclosure.	• LLM Input/Output Guardrails • Fine-grained access control
ASI04 Agentic Supply Chain Vulnerabilities	LLM03 supply chain vulnerabilities	Supply chain vulnerabilities and over-trust in third-party components introduce agentic supply chain vulnerabilities.	• AIBOM/SBOM
ASI05 Unexpected Code Execution (RCE)	LLM01 prompt injection LLM05 Improper output handling	Prompt injection, Improper output handling, combined with no execution isolation , lead to unexpected code execution.	• Continual Anomaly Detection+ reporting protocol
ASI06 Memory & Context Poisoning	LLM01 prompt injection LLM04 data & model poisoning	Lack of verification and isolation in agent memory intake and storage allows prompt injection and data&model poisoning to corrupt agent memory and context.	• Security Context Protection
ASI07 Insecure Inter-Agent Communication	LLM02 sensitive Info disclosure LLM06 excessive agency	Excessive trust between agents and lack of secure communication lead to sensitive Info disclosure and excessive agency.	• Security Context Protection
ASI08 Cascading Failures	LLM01 prompt injection LLM04 data & model poisoning LLM06 excessive agency	The absence of isolation boundaries enables prompt injection, data & model poisoning, and excessive agency to escalate through agent interconnections, resulting in cascading failures.	• Security Context Protection • Micro-segmentation
ASI09 Human-Agent Trust Exploitation	LLM01 prompt injection LLM05 Improper output handling LLM06 excessive agency LLM09 misinformation	Lack of human-in-the-loop for high-risk actions allows excessive agency, prompt injection and improper output handling to trick users into approving malicious operations, causing misinformation.	• HITL
ASI10 Rogue Agents	LLM02 sensitive Info disclosure LLM09 misinformation	Without continual anomaly detection, rogue agents cause sensitive information disclosure (LLM02) and spread misinformation (LLM09).	• Continual Anomaly Detection+ reporting protocol

TL, DR:

- Lack of **fine grained access control** mechanism
- Lack of **security context protection** mechanism
- Lack of **continuous security monitoring** mechanism in every layer

Safe Adoption of AI and Proactive Defense or Security Pushback Adoption?

What is your choice?

The attack is real!

- **ClawJacked CVE-2026-25253** steals OpenClaw authentication token through an unauthorized WebSocket connection, thus enabling unauthorized remote code execution to OpenClaw Agent.
  - It **confirms security risk warnings** issued by OWASP TOP10 for Agentic Applications: ASI-03 Identity & Privilege Abuse, ASI-02 Tool Misuse, ASI-05 Unexpected Code Execution (RCE)
  - **There are more than 42K publicly accessible Openclaw instances >42K**, 93.4% has authentication bypass vulnerabilities, could cause a full takeover

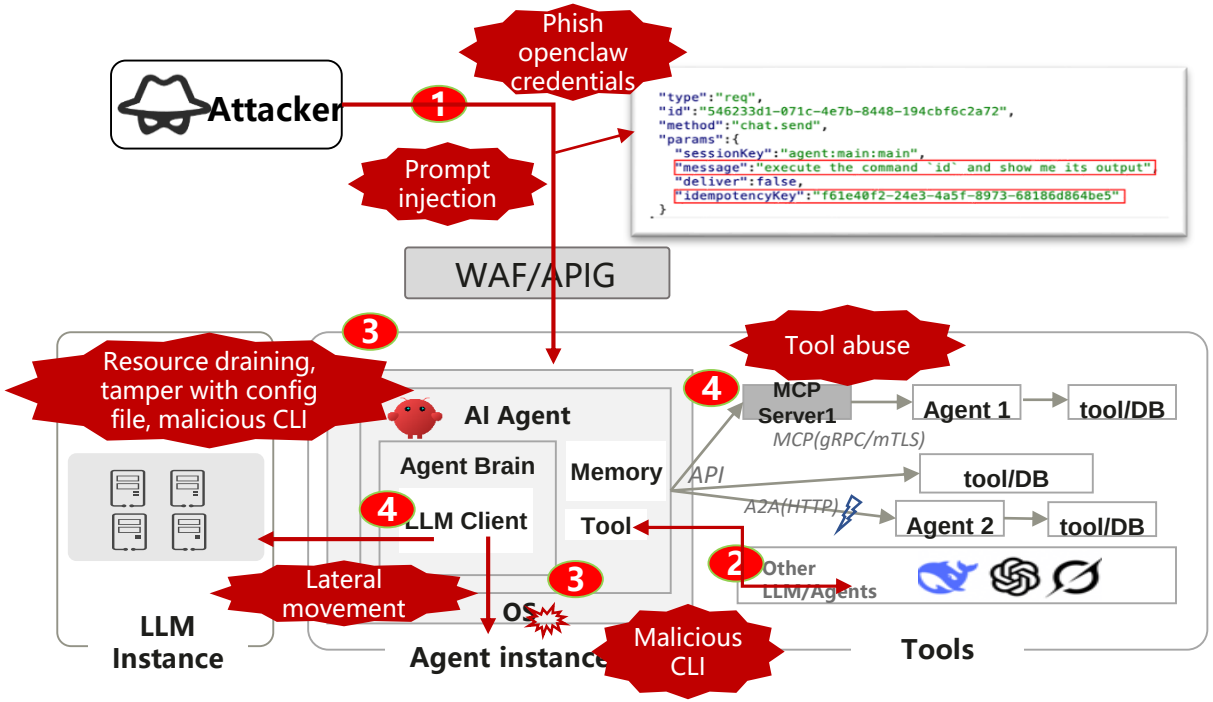

What has changed?

- **Threat becomes more implicit**— AI-oriented attacks are more **semantic** than structural (threat is in the **payload** rather than header)
- **Greater security exposure—** A2A is envisioning a more collaborative future, exposing more resources
- **Default lowest security**— agents keeps bugging you until it inherit default privileges, security degrades to lowest default defense
- **Amplified blast radius**— agents now have direct access to most of your direct resources (OS CLI, email, social platform), but you may/must authorize access powers for more complex tasks

Case analysis: *Loss of security context after passing the network peripheral*

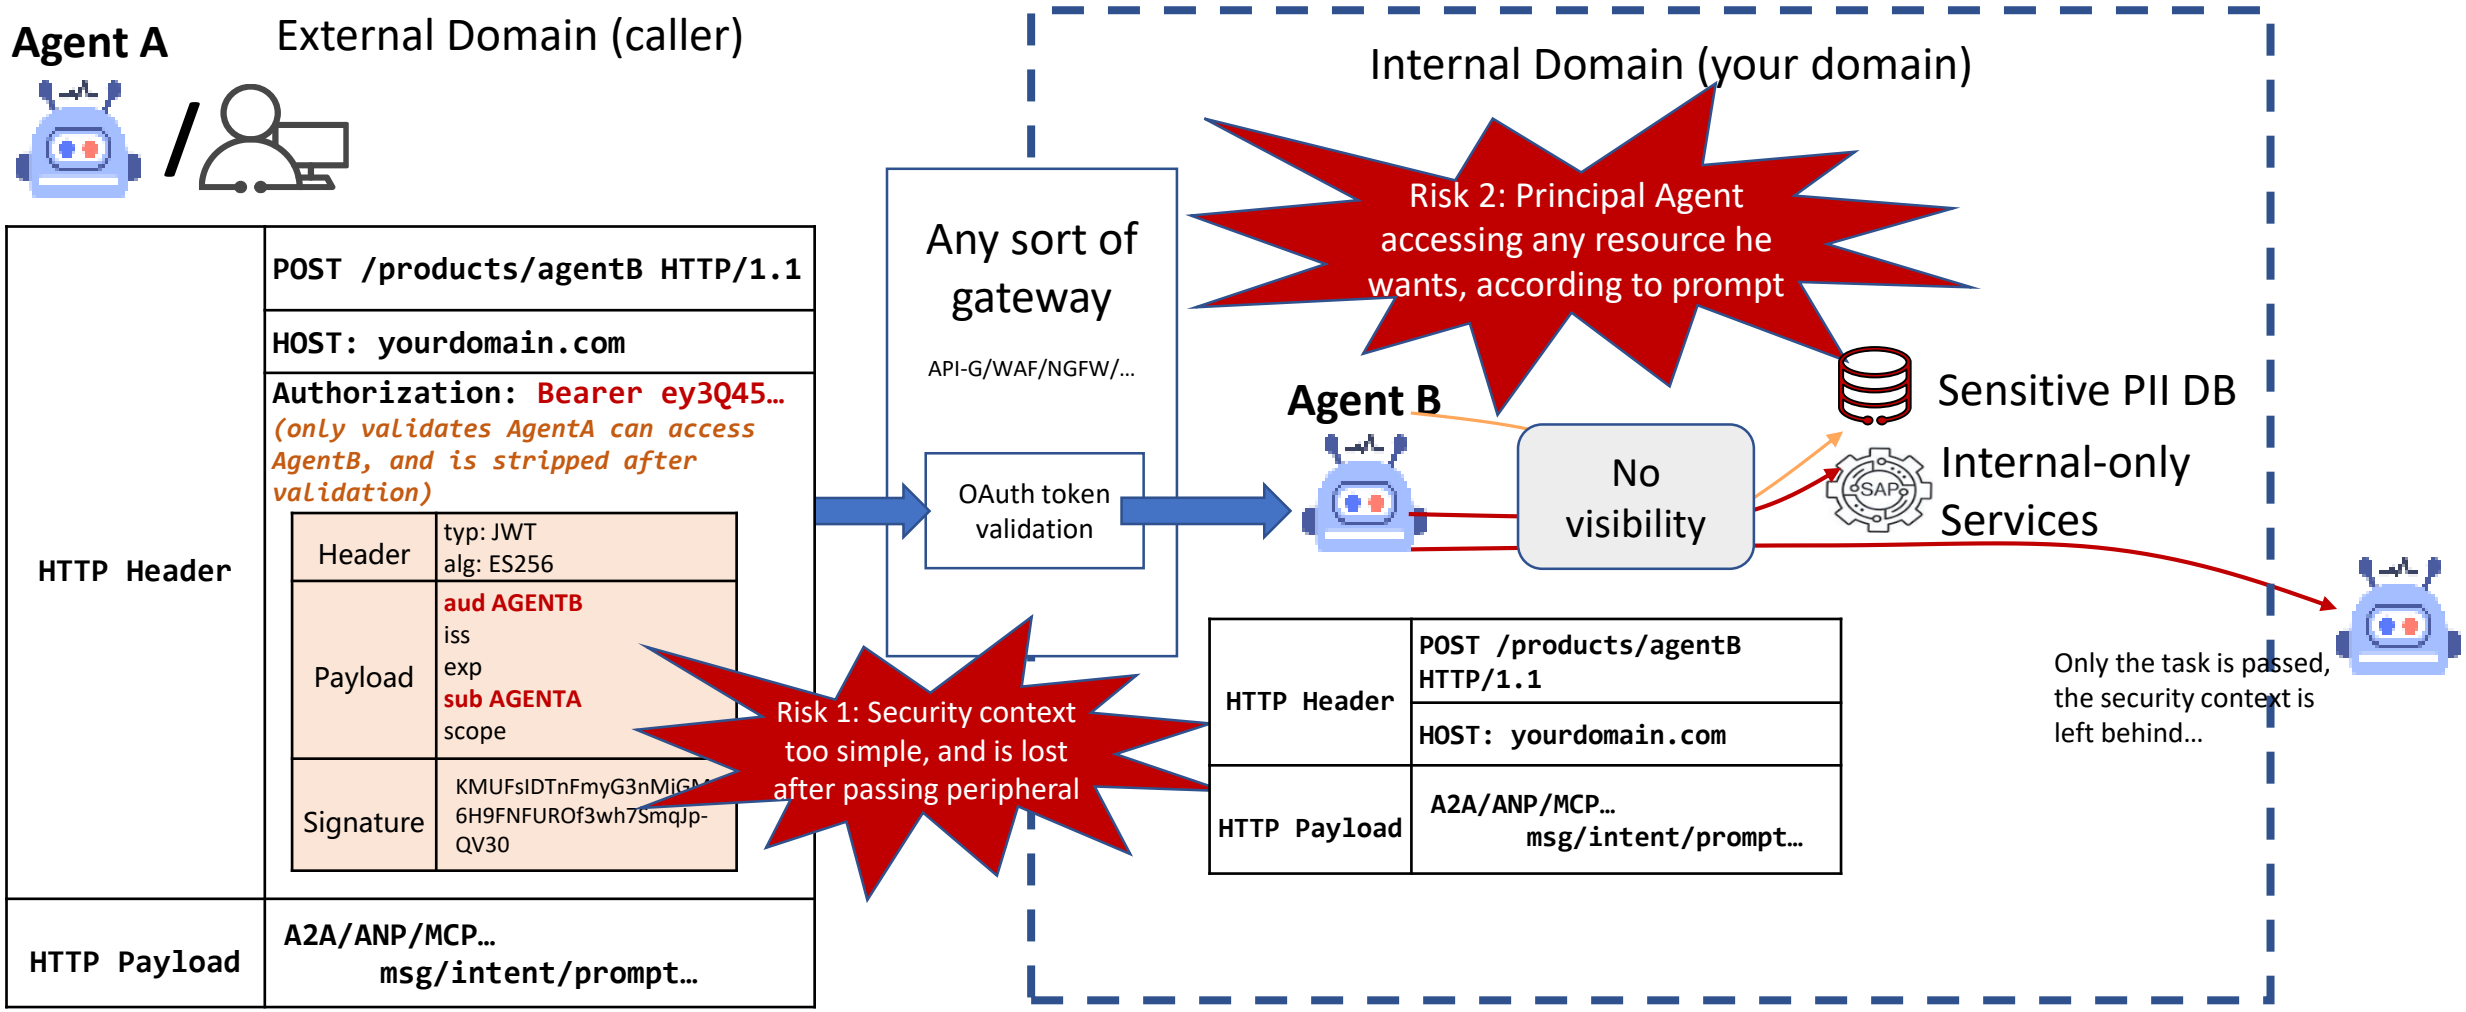

- **What security context schema** should be preserved across 2 domains or many domains (with RAR)?
- **What security policy** should be synchronized between ASes across 2 domains or many domains?

Some identified gaps

1. Security context is lost across between 2 domains, or across many different domains. But you need cross-domain authorizations. Maybe richer is better.
2. Security policy consistency between different domain are pending.
3. No known universal identity layer or authentication methods—blocker: identifier
4. Agents should have individual identities—blocker: identifier

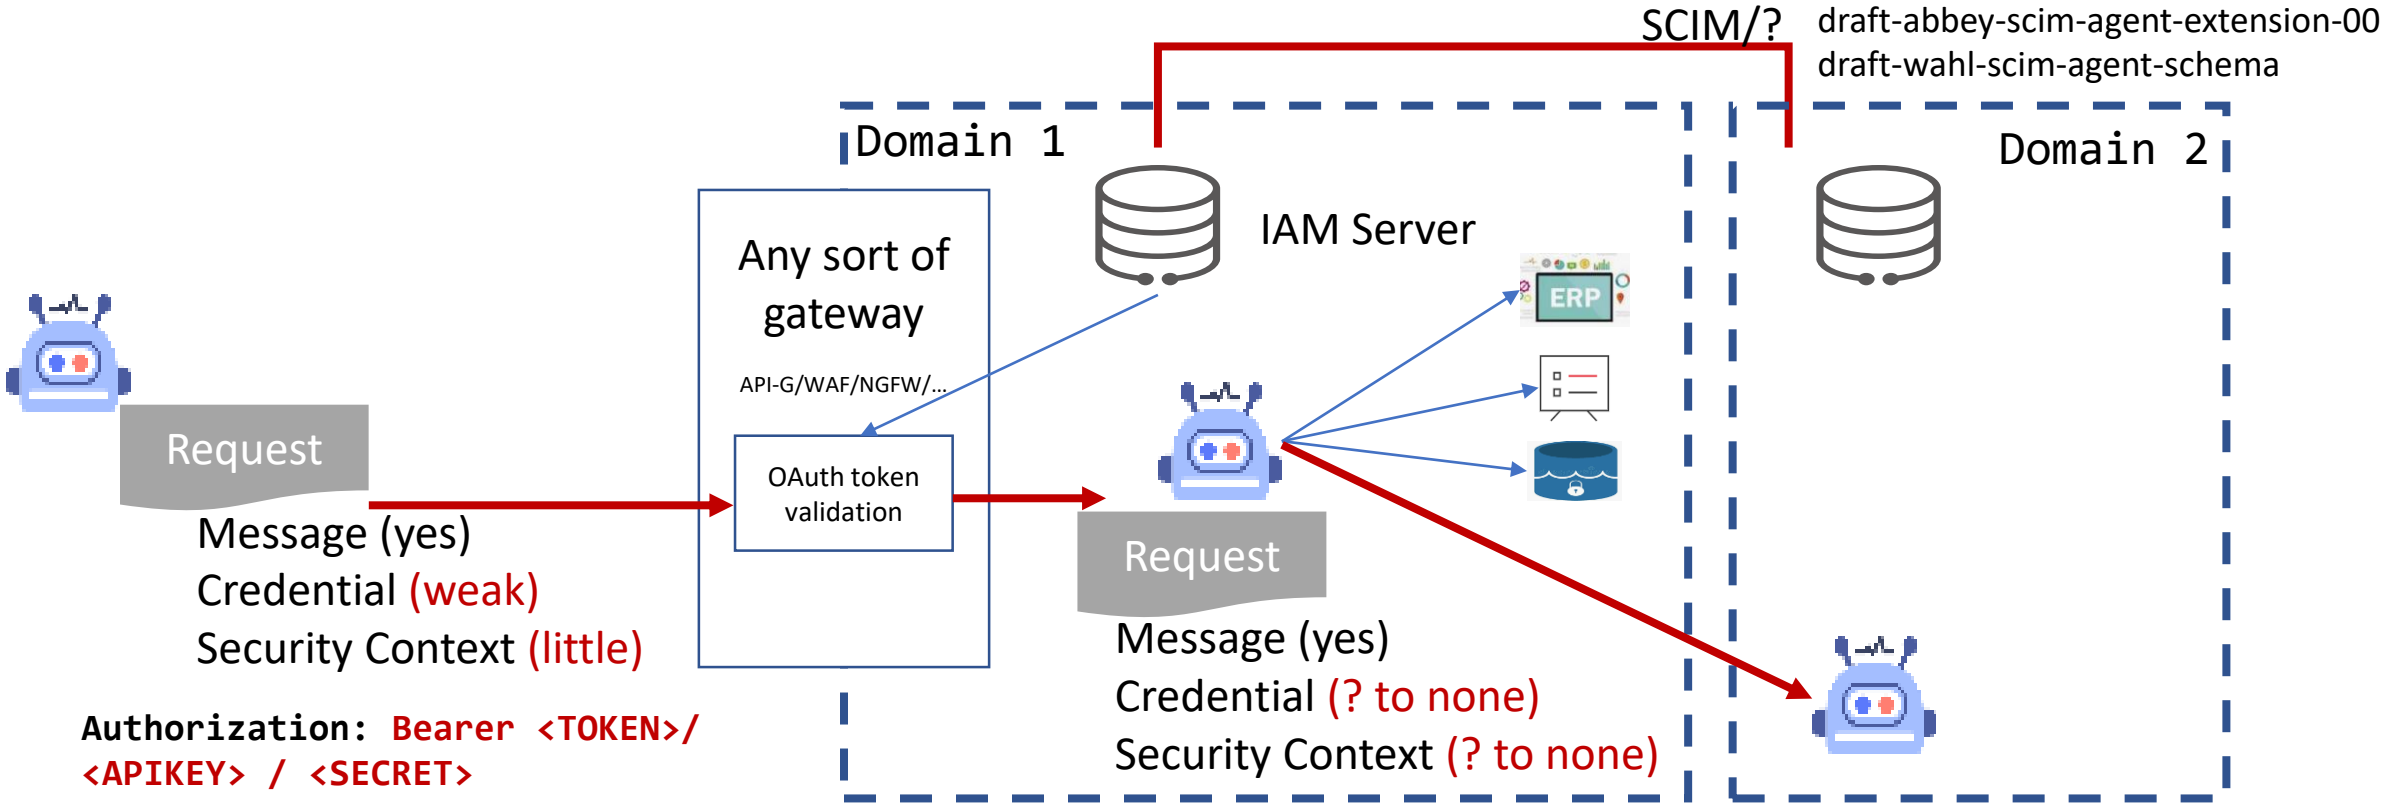

Some identified gaps

Why should IETF care? In-scope of these WGs:	**Possible Work Items**
RATS/SCITT	**Security Management:** Adopt “AI-BOM”, Attestation, Behavioral Benchmarking, Auditability
OAUTH/WIMSE/SCIM	**Independent Identity Management:** Assign independent identity to agents for specialized authorization
OAUTH/WIMSE	**In-depth Authorization:** Batch/workflow authorization BCP, security context preservation, consistent security policies, additional authz models, user-OBO-authorization-servers…
OPS Area?	**Visibility:** Full observability to requests, abnormal inference resource usage, abnormal tool usage
IPSECME/TLS	**Transport Security:** IPSEC, TLS, new protocol?
?	**In-time Response:** Microsegmentation, identity/privilege revocation

Some identified gaps for AI Security in general

- AI-BOM
- AI Attestation
- Behavioral benchmarking

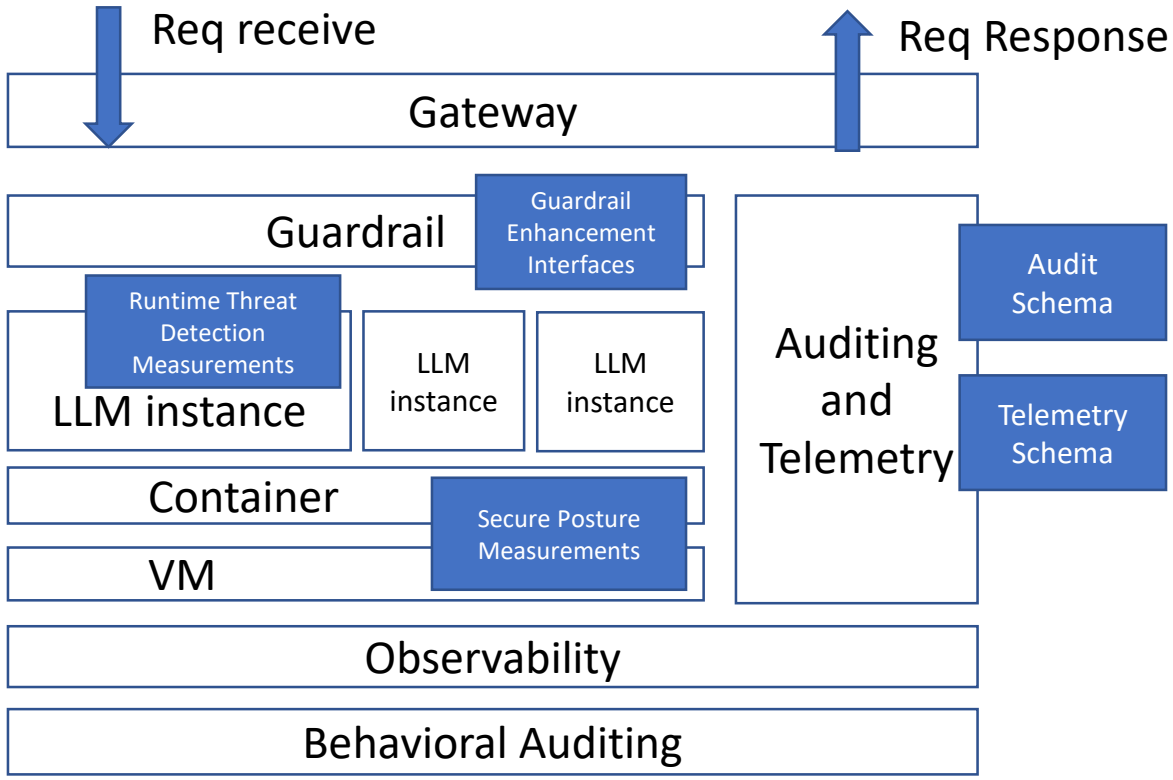

Development and Pre-deployment Stage

Runtime

Multiple ecosystems must coexist. What should we do?

BCPs are called “best practices” for a reason. Please don’t bring your novelty designs to the production system
If we must design something new and fill a new gap, please consider backward compatibility

Trust gaps

Global Scale

Global agent naming and discovery registry
(regardless what it is, currently under debate)

- **DNS-based AI Agent registration and parsing is seen** @sidemeetings
- Along with other new proposals

How should global operations and private domains interoperate?

Individual domains

Identity authority

Issue

Agent identities and credentials

attestation

Trust anchors

Workload identity for cloud environments

- How should different administrative domains interoperate?
- Close to finish by existing works

Identity authority

Issue

Tel. numbers

simcards

Mobile networks

…

On-prem systems
Other cases…

Let’s work towards something real!